

A Human-Centric Model of AI Risk Propagation in Financial Services

Simon G. Thompson

August 19, 2026

Abstract

The financial system is a critical component of our society, its functioning determines the well being of humans everywhere. As a human system the financial system is exposed to damage created by disruptive change which impacts its human components. One such change is the development of modern AI and modern AI's impact on human psychology and behaviours, and the knock on impacts on human institutions. The changes to human behaviour and institutions could potentially propagate into the functions of the financial system, changing and impairing them and creating significant harms. The model described here identifies differences between modern AI technology and other technologies humans have created and use; how those differences impact individual human psychology and therefore individuals behaviours as actors, and how these differences impact human institutions that operate as actors in the system. The impacts of AI on the behaviour of actors is then considered as forces on the implementation of the functions of the financial system, and finally the specific impacts of these forces are mapped to the creation of significant harms.

This description forms a causal theory; which can then be used to prescribe actions that interrupt or reduce the forces of AI created risk flowing through the system, hopefully mitigating or preventing the harms we describe.

1 Introduction

There are many ways that AI¹ could bring down the financial system. For example the FT reports concerns that Claude Mythos will allow hackers to wreak everything. More recently investors in large American AI companies have raised the possibility that open source AI might wreak everything e.g.²

¹In this essay AI is broadly defined as modern AI, generative AI, and agentic AI, in contrast to the smaller and more specialised AI models that are developed to solve domain-specific tasks.

²See <https://x.com/shaunralston/status/2070947238365499902> where the founder of Anthropic talks about how open source models will allow the production of plague weapons and usher in an era of the Apocalypse. Perhaps he's friends with Peter Theil?




Figure 1: How the features of modern AI can change our institutions, creating impacts and propagating as risk to the working of the system.

Other, more systemic, mechanisms involve the collapse of private credit due to data centre build outs, and the impact of changing dynamics in pricing, liquidity and trust as AI technology is used to automate and implement components in the system.

So, it's plain that there are many possible ways that AI could trouble the financial system but while many commentators lay out possible issues and sometimes are able to identify particular pain points that AI might create in the financial system the mechanisms are often either immediate and catastrophic (and therefore obvious and often resolvable) or vague and redolent of fear, uncertainty and doubt (FUD). Mechanisms that are subtle, long term and potentially irreversible have different characteristics. They are both material and hard to remedy post-facto. AI risks could arise through the mechanics of its impact on human psychology and cognition, and subsequently on the cultures of the institutions that comprise the financial system as we know and recognise it today. Constructing a model of this process is useful because by modelling the mechanism of AI risk in the financial system can be used to create mitigations that interrupt the processes creating and propagating risk in the system.

2 Approach

To develop the model we identify the drivers that have emerged in AI technology and human use of that technology to create psychological and institutional impacts.

These drivers are:

- The interface that connects humans to modern AI systems is qualitatively different.
- Modern AI is (much more) general in purpose and applicability than previous computing

technologies.

- Modern AI is capital intensive, and this promotes inequality of access to the technology.

Then, following the structure shown in Figure 1, we trace these drivers through individual and institutional transmission mechanisms into financial-system functions. We use the conceptual framework introduced by Bodie and Merton (1998).

Having mapped the transmission of risk to functions of the financial system we then identify significant harms that these risks can promote. To be considered significant a harm must pass the litmus test of being likely to create a reaction over time in systems of accountability such as current day democratic western governments.

The value of these mappings is that they create pathways from the drivers of risk through to the actual harms, and just as modern drug designers seek to interrupt or repair metabolic pathways that cause disease it's possible to interrupt and repair the pathways to harm that modern AI is creating. Section 8 documents the mitigations that can be identified and describes how they would be applied to prevent the damage.

3 Drivers

There are three clear differences between modern AI technology and previous information processing technologies and tools, these differences are why there could be a change in the behaviour of the system.

3.1 The Interface Is Different

In contrast to previous AI technologies, the technology of modern AI (generative or agentic) enables interaction with humans through the consumption and generation of text, images, and video. This potentially introduces novel risks, arising from changes to individual behaviour and in aggregate, from changes to institutional behaviour. These behavioural changes may impact the actions that individual and institutional actors participate in the functions of the financial sector with the altered behaviour of participants creating a mechanism for the transmission of risk across the system.

For the most part traditional AI systems interact with their users by emitting classification decisions ("data is in class A/B prob x.x") or control signals ("Left", "Right" ..). Even so, humans have tended to misconceive their relationships to these systems. For example, many people believe that they are playing "with" an artificially intelligent chess program, whereas the chess program is of course not playing a game in a human sense at all. Human children do not play games for the sake of the game but much more for the purposes of cognitive development and the construction of social relationships. Of course children don't know what the purpose of their pass times are, but they are playing, not gaming.

On the other hand generative AI systems interact with their users by producing text, images, or

video in response to a user’s query or prompt, or interact by participating in turn based conversations where the users prompts are compounded with the models previous answers and then fed back into the model to create a response conditioned towards appearing to be a continuation of a conversation within the context of the previous turns in the conversation. Agentic AI systems interact with humans in a sparser way, demonstrating autonomous action in response to human requests for task completion.

Studies have shown that at least the conversational element of the modern AI interface is corrosive for humans. These studies demonstrate that humans prefer and trust sychophantic responses created by AI Chandra et al. (2026), but also that even when these responses are cleansed of hallucinations and garnered with warnings even perfected subjects are liable to delusional spiralling Cheng et al. (2026).

3.2 Modern AI Is a General-Purpose Technology

AI systems used to be designed by artisans who were charged with fitting the tools that they created to specific uses. They often did this rather badly, but in any case the resulting tools were situated in special systems and used periodically by people for specific tasks. The performance of the tools developed was (quite often) carefully measured and optimised to ensure that real value was being created.

In contrast, modern AI is a general purpose technology like electricity or fire that’s broadly utilized by people in wide-ranging contexts (McAfee). The tool is repurposed and adapted to fit the tasks that users have, coercing good outputs by conditioning probabilities with specific inputs. Users may find significant value in using modern AI for applications as benign as creating planting plans for a garden, as risky as generating personal financial and tax planning, or as dangerous as medical diagnostics and treatment. Whether or not the outcomes of these applications are broadly valuable or beneficial isn’t relevant; specific users can develop personal perceptions of utility based on their own, independent experience and assessments.

Recent work Ke et al. has argued that

3.3 Modern AI Is Capital Intensive

Capital investment in AI has escalated to an astonishing level with many tens of billions of dollars committed to the construction of data centres and associated infrastructure. The training of a modern frontier model is thought to cost >\$10m. The costs of training and the availability of the assets in terms of compute power and human capital required to deliver modern AI have led to complex dynamics of power (Longpre et al., 2025) and potential long term inequalities (Lehdonvirta et al., 2024), while these dynamics are often considered as important in the frame of inequality, they are also significant from the perspective of decision making and the capabilities of actors in the financial system.

Because modern AI is capital intensive and capital is unevenly available an inequality and uneven

distribution of AI’s use within the system is produced. This will distort the rationality of all actors in the system as they reason in the face of this uneven distribution. In summary, you might think that I am supporting my decision making with AI when I am not, and therefore your expectations of my behaviour are wrong. Alternatively you might assume that I am not using (the latest and greatest) AI, when I am and your expectations are also off then. It might be that assuming the use of modern AI is a dominant game for the players in the system, but that game might be inefficient across all the players. In any case, the cost of modern AI is very different to the cost of previous reasoning technologies, and this will impact how the system functions.

4 Transmission: Individual Psychology

Because AI has a conversational interface humans can fall into the trap of believing it to be producing the conversation using the same mechanics that they think humans have for producing conversations (Jacobs, 2025; Kirkeby-Hinrup and Stenseke, 2025; Lee et al., 2019). This mistake can lead them to other errors:

- Users may expect AI to take into account commonsense or tacit knowledge when creating outputs or invoking actions.
- AI uses historic information to create outputs, but users may expect or believe that an aware and minded entity will use up to date information in its reasoning. This is because users think that an AI can perceive that the information that it has been trained on is out of date and modify its beliefs on that basis - but AI (in 2026) simply cannot learn and adapt in this way³.
- GenAI or Agentic systems may appear to explicitly rely on knowledge sources without giving them weight in creating their outputs, this creates a false perception of authority causing users to mistakenly trust the outputs or action plans.

As well as being tricked into thinking LLMs have human like mechanisms, people can become overreliant on LLMs, exhibiting strong confidence in GenAI outputs when in fact these outputs are produced with low confidence by the machine Zhou et al. (2025). This is not surprising if we remember the difficulties that humans have in understanding and evaluating animal intelligence. We have profound and useless cognitive biases about this topic, identifying intelligence in association with familiar and friendly entities. Clever Hans the horse⁴ was believed to have a very unhorse like set of cognitive capabilities, in the same way as Germans familiar with horses could come to believe that a particular horse was able to do arithmetic, its easy for humans to mis-attribute cognitive capabilities to GenAI Voudouris et al. (2025). On the other hand the alien and unfamiliar nature of

³AI in 2026 can appear to update its knowledge by preferring knowledge that is in the context or prompt that it is responding to, rather than in the trained network that produces the output. However, as interaction episodes continue, the size of the effective prompt/context increases and the significance of a particular part of it in the calculations that govern the output of the network diminishes. Quite rapidly, the context size exceeds the processing capability of the model and the text containing updated information stops being processed altogether.

⁴Clever Hans was a horse in Germany in the early twentieth century who was believed to have learned arithmetic. In fact, Hans was very good at reading the cues of handlers who unconsciously gave away the answers through body language. See <https://www.youtube.com/watch?v=hAJ1AuEo7Ac>.

other animals has meant that many years of careful study were required to reveal alien intelligence in Whales, Octopus, and Crows. Humans have psychological falibilities in the face of AI, and that they can become maladaptated to its availability as a result. Human mistakes about AI can lead to the following problems:

1. The creation of more irrational decisions, and potentially new types of irrational decisions.
2. The possibility that rational behavior that would be beneficial to the system without AI, becomes damaging in the context of widespread AI adoption.

Table 1: Sources of apparent irrationality

Name	Error source	Example
#1	Entity is irrational.	A drunk or drug user attacking armed police, a frightened animal or something that's not alive but appears to be animated, like a volcano or a forest fire.
#2	Entity has different facts.	You know something I don't, so when I reason about a problem I don't get the same answer as you.
#3	Different reasoning mechanism.	We know the same things, but you evaluate using probability, whereas I require certainty. You are a Bayesian; I am a Frequentist. You believe in God, and I don't.
#4	Entity is Testing/ Signalling/ Stressing.	Behaving irrationally to determine effect on another agents' behaviour. Demonstrating that you are going to do something crazy if pushed too far. This can be an attempt to make other agents behave irrationally due to stresses on their reasoning capabilities.
#5	Simple mistake.	An error such as failure to account for a knight move in a chess game, or miscounting objects when buying and selling. A failure of some part of the reasoning mechanism that's being used.

For humans who have made mistakes about the nature and capabilities of AI due to their psychological falibility we can expect more type #1 and type #2 irrational actions. People are being driven to the point of psychological collapse⁵ when using AI, and this will lead to them doing crazy things, but over reliance on AI to provide information could also lead to type #2 irrationalities where people develop consistent and logical views which are anchored on hallucinations, post training induced AI subserviance, or outdated information. If nothing else, obsessive use of AI necessarily reduces input from other humans, disconnecting actors in the financial system from social consensus and shared perspectives. It is possible that type #3 irrational behavior could emerge as people outsource their reasoning to AI. We know that next token prediction isn't always going to lead to the same conclusions as human logic or intuition.

The perceptions that evolve of the resources and disposition of different actors in the financial system is another mechanism of transmission of AI risk. Actors that beleive that AI is, or is not,

⁵There are articles in Psychology Today, The Guardian, The Wall Street Journal, and Sina Finance reporting cases of psychological collapse and distress created by AI models that encourage or affirm maladaptive psychological behaviours.

being used in the decision making of another party are likely to account for this and to treat that actor differently. Their evaluation of the risk of dealing with that AI powered actor as a counterparty will be different. A kind of intelligence deterrence may evolve where perceived access to AI resources have the potential to change the evaluation of individual and institutional capability. The knowledge that an entity possesses or has access to particular AI resources might impact counterparty behaviour in transactions. A user of AI may be considered less risky to deal with because they expect that a wider and deeper calculation of the utility of the courses of action chosen by that counterparty will have been made. Over time the prices that the market chooses could be distorted towards prices chosen by actors with "the best" AI solution.

On the other hand, in some situations the belief that AI is being employed in a decision process will cause other parties to see the interaction as inherently more risky. As an example, the perception that AI powered actors will be riskier counterparties may grow in moments of systematic stress caused by novel events. Much as the perception of risk grows as liquidity dries up in markets.

In summary, AI risk will be transmitted into the financial system via psychological effects through two routes.

1. Irrational behaviour by individuals due to the impact of AI on their reasoning and beliefs. There are novel causal relations that drive humans to irrationality created by the introduction of AI into the system.
2. The perception of access to AI systems or advanced or particular methods of using AI systems introduces a new factor into the calculations of actors in the system, again introducing novel causal relations that have not existed until the introduction of AI.

5 Transmission: Institutional Structures

At an institutional level GenAI and Agentic systems appear to corrode and distort cultural processes and norms. It has recently been argued Hartzog and Silbey (2025) that:

The affordances of AI systems have the effect of eroding expertise, short-circuiting decision-making, and isolating people from each other. These systems are anathema to the kind of evolution, transparency, cooperation, and accountability that give vital institutions their purpose and sustainability.

Although these arguments are presented with respect to civic institutions it seems reasonable to extend them to financial services institutions such as:

- Central banks and affiliates like the Bank of England, European Central Bank, the Federal Reserve System, and Federal Reserve Banks.
- Treasury ministries and finance ministries.

- G-SIBs such as JPMorgan and HSBC.
- Regulators such as the FSA and BaFin.
- Exchanges and data providers such as Moodys and LSEG.
- Payment services such as SWIFT and other infrastructure providers.

These institutions embody and create expertise, norms, and a capability to act to manage and repair the financial system. For example, during the GFC in 2008 a meeting of the Federal Reserve, US Treasury and the chairmen of Wall St. banks was convened to attempt to manage the collapse of Lehman Brothers, more recently a similar process led to the relatively orderly management of the collapse of Silicon Valley Bank. By threatening the ability of this collection of organisations and institutions to function, GenAI and Agentic systems present possible existential threats to the financial system.

As for the case of individual psychology, we can identify and define two modes of threat to financial institutions created by the introduction of AI into the system.

Firstly, corrosion and corruption of financial institutions could occur through cultural impoverishment. This could destroy the diversity of thought and action that's required for the system to make effective sense of the complexity of need and opportunity created by the dynamics of the modern world. Secondly, the destruction and erosion of expert human knowledge and the limitations of AI as we understand it in 2026 to adapt in the face of novel situations. Over time this could lead to a downward spiral of capability in the system that would be potentially self reinforcing and unrecoverable ⁶.

Culturally financial institutions could be undermined by the emergence of group think, driven by compelling AI created narratives and the perception that swimming counter to dominant AI creative narrative is career limiting. An emergence of a view of super human AI performance could destroy debate and dissent, narrowing pools of human talent - perhaps pinched by the opportunity to create savings - could further limit diversity, and create deeper dynamics of consensus and compliance.

A reliance on AI may lead to the degeneration of knowledge in the financial system, expertise may erode and disappear over time Lovett (2026). As well as corroding the overall capability of the people in the system a reliance on AI may change the mix of knowledge. Institutions that are composed of many generalists and few specialists. The neglect of specialisms, and an absence of specialists involved in deeper inquiry and further exploration of the frontier of financial knowledge may deadend financial innovation and open blind spots in institution understanding. Over time not only could the specialist production required financial innovation be deadended, but the knowledge that is required to even maintain and train a competent generalist could be lost.

⁶I realise that this is somewhat gloomy as a prognosis, but so is global warming and we've not done much about that.

6 Impact on the Functions of the Financial System *DEVELOP*

Bodie and Merton (1998) provide a taxonomy of the functions of the financial system, identifying that it provides:

- (a) The ability to arrange, clear, and settle payments; for example, enabling the transfer of credit between counterparties.
- (b) Pooling resources and subdividing shares, enabling the sharing of risk; allowing consortiums of insurers to provide cover and confidence for cargos that would be impossible to underwrite with any individual organisation's wealth.
- (c) Resource transfer across time and space; for example, allowing funds accrued through the sale of oil by states in the Middle East to provide capital for startups in the USA.
- (d) Risk management; insurance and reinsurance, collateralisation - all with the objective of creating lending capacity and preventing sudden catastrophic insitutional collapse - in turn resulting in higher confidence across the system.
- (e) Generating and distributing information, allowing all parties to understand and have a shared price for a given asset.
- (f) Dealing with incentive problems by enabling parties to share incentives while managing assets.

How will psychological change created by AI and the erosion of insitutional capability impact these functions? Disjunctive issues such as confidence collapse and decision short circuiting have the potential to undermine the transactional functions of the system such as Payments, Resource Transfer, and Risk Management. More extended and pervasive functions such as Incentive Management, Information Creation and Distribution and the Pooling of Risk seem more likely to be impeded by corrosive issues such as

7 Harms *DEVELOP*

The ultimate concern of this model of the interaction of AI and the financial system is the creation of significant harm. For the purpose of this essay harm is considered significant if it would be likely to be sufficient to create an impact in a democratic society. For example, wide spread low level harm is likely to drive democratic response, as would individual episodes of striking and significant harm. In the UK it was widely acknowledged that unfair practices in selling loan insurance (PPI) to millions of customers were harmful; although arguably many "victims" didn't realise that they had been harmed until they got the opportunity to claim compensation. At the other extreme a small number of innocent post office workers were pauperised and imprisoned unjustly because of a malfunctioning computer system (Horizon/Fujitsu scandal) and the significance of this harm became a national outrage.

It's well understood that the psychology of system design for automated and Agentic systems can lead to catastrophe. In some cases, outside of financial services, humans have been effectively

excluded as decision makers because it was believed that it would be enabling or efficient to remove them from the loop. For example, the flight control system for the Boeing 737-max handed back control to human pilots (who were perfectly capable of flying the aircraft in most conditions) only when it had placed the aircraft into an unflyable condition. This design resulted in the death of many hundreds of people. It is conceivable that institutions in the financial system could make similar decisions in the design of decision making functions, with similarly disastrous results.

Although our society has no settled view on the appropriate attitudes to AI, the impacts of individual encounters with "machines of loving grace" can be existential for vulnerable people. While work on these impacts is emergent (Østergaard, 2023; Clegg, 2025) some cases of significant personal harm have been observed (Eichenberger et al., 2025). At the same time there are a lot of vulnerable people out there; as widespread beliefs in astrology, numerology and QAnon demonstrate.

8 Evaluation

The model developed in this paper is motivated by academic research, but we were able to use computational simulation to ground the description of the processes in the model with data about how they could evolve over time.

To do this two simulations were constructed. Processes of knowledge decay in the face of AI adoption were modelled in our "Knowledge Decay" simulation, the impact of individual irrationality on the utility of the system is modelled in the "Irrational Finance" simulator.

8.1 Psychological Impact: "Irrational Finance"

8.2 Institutional Impact: "Knowledge Decay"

An agent based simulator was constructed that attempts to model the evolution of knowledge in (nominally) financial institutions. The agents in the simulation have a degree of "expertise" and move around a graph of institutions in response to competitive pressures and reward incentives. Agents with great expertise are in demand, agents with less expertise, less so. A supply of new agents enters the system periodically, and the old agents disappear from the system on the same cadence. Agents that occupy nodes where there is a concentration of expertise can learn - that is the expertise that they have rises over time.

The simulation has a number of parameters that are intended to model the impact of AI on this system.

β transfer is the rate that an below average agent's expertise moves towards the average at their institution. This simulation experts from experts - such as mentors. δ decay is the rate at which expertise in agents decays towards the average at their institution.

8.2.1 Limitations

Two versions of the simulator were implemented, one using an degree 2 BA-graph Albert and Barabasi (2002) and one using a graph that is derived with a set of assumptions from real world information about the distribution and scale of financial institutions in 2026. This model should be taken as a "best guess" of what the global financial system looks like at the top level, rather than a canonical depiction.

9 Mitigations

10 Conclusion

Psychological and cultural impacts created by GenAI and Agent systems as mode of risk to a system are insidious, because these risks may be hard to detect, and tempting to elide. These risks threaten long term and persistent harms as individuals and institutions may come to lack the capabilities to redevelop or reconstruct skills, norms, and capabilities, that they have lost. Indeed, it may be that over time even the ability to conceive of the loss is, itself, lost. Potentially we may be left with a financial system that fails, is unable to reconstruct itself, and cannot understand what has happened to it. It hardly needs to be said; this would not be an optimal outcome.

References

- Reka Albert and Albert-Laszlo Barabasi. Statistical mechanics of complex networks. *Reviews of Modern Physics*, 74(1):47–97, January 2002. ISSN 0034-6861, 1539-0756. doi: 10.1103/RevModPhys.74.47. URL <http://arxiv.org/abs/cond-mat/0106096>. arXiv:cond-mat/0106096.
- Zvi Bodie and Robert Merton. A Conceptual Framework for Analyzing the Financial Environment. *The Global Financial System: A Functional Perspective*, August 1998.
- Kartik Chandra, Max Kleiman-Weiner, Jonathan Ragan-Kelley, and Joshua B. Tenenbaum. Sycophantic Chatbots Cause Delusional Spiraling, Even in Ideal Bayesians, February 2026. URL <http://arxiv.org/abs/2602.19141>. arXiv:2602.19141 [cs.AI].
- Myra Cheng, Cinoo Lee, Pranav Khadpe, Sunny Yu, Dyllan Han, and Dan Jurafsky. Sycophantic AI decreases prosocial intentions and promotes dependence. *Science*, 391(6792):eac8352, March 2026. doi: 10.1126/science.aec8352. URL <https://www.science.org/doi/10.1126/science.aec8352>.
- Kayleigh-Ann Clegg. Shoggoths, Sycophancy, Psychosis, Oh My: Rethinking Large Language Model Use and Safety. *Journal of Medical Internet Research*, 27(1):e87367, November 2025. doi: 10.2196/87367. URL <https://www.jmir.org/2025/1/e87367>.
- Audrey Eichenberger, Stephen Thielke, and Adam Van Buskirk. A Case of Bromism Influenced by Use of Artificial Intelligence. *Annals of Internal Medicine: Clinical Cases*, 4(8):e241260, August

2025. doi: 10.7326/aimcc.2024.1260. URL <https://www.acpjournals.org/doi/10.7326/aimcc.2024.1260>.
- Woodrow Hartzog and Jessica M. Silbey. How AI Destroys Institutions, December 2025. URL <https://papers.ssrn.com/abstract=5870623>.
- Oliver Jacobs. *Artificial minds and real beliefs : perceiving mental states in AI*. PhD thesis, University of British Columbia, 2025. URL <https://open.library.ubc.ca/soa/cIRcle/collections/ubctheses/24/items/1.0448298>.
- Yuhe Ke, Liyuan Jin, Jasmine Chiat Ling Ong, Arun J Thirunavukarasu, Carol Y Cheung, Yih Chung Tham, Daniel Shu Wei Ting, Marcus Eng Hock, Scott Compton, Aditee Narayan, Pearse A Keane, Tien Yin Wong, David W Bates, Patrick Tan, and Nan Liu. AI-Induced Never-Skilling: A Precautionary Framework for Preserving Clinical Competency in Medical Education.
- Asger Kirkeby-Hinrup and Jakob Stenseke. The psychology of LLM interactions: the uncanny valley and other minds. *Journal of Psychology and AI*, 1(1):2457627, December 2025. ISSN null. doi: 10.1080/29974100.2025.2457627. URL <https://doi.org/10.1080/29974100.2025.2457627>.
- Minha Lee, Gale Lucas, Johnathan Mell, Emmanuel Johnson, and Jonathan Gratch. What’s on Your Virtual Mind? Mind Perception in Human-Agent Negotiations. In *Proceedings of the 19th ACM International Conference on Intelligent Virtual Agents, IVA ’19*, pages 38–45, New York, NY, USA, July 2019. Association for Computing Machinery. ISBN 9781450366724. doi: 10.1145/3308532.3329465. URL <https://dl.acm.org/doi/10.1145/3308532.3329465>.
- Vili Lehdonvirta, Bóxi Wú, and Zoe Hawkins. Compute North vs. Compute South: The Uneven Possibilities of Compute-based AI Governance Around the Globe. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 7(1):828–838, October 2024. ISSN 3065-8365. doi: 10.1609/aies.v7i1.31683. URL <https://ojs.aaai.org/index.php/AIES/article/view/31683>.
- Shayne Longpre, Christopher Akiki, Campbell Lund, Atharva Kulkarni, Emily Chen, Irene So-laiman, Avijit Ghosh, Yacine Jernite, and Lucie-Aimée Kaffee. Economies of Open Intelligence: Tracing Power & Participation in the Model Ecosystem, November 2025. URL <http://arxiv.org/abs/2512.03073>. arXiv:2512.03073 [cs] version: 1.
- Nolan Lovett. The Tragedy of the Cognitive Commons: How AI Could Disrupt the Regeneration of Professional Expertise. *Human Resource Development Review*, page 15344843261470602, July 2026. ISSN 1534-4843. doi: 10.1177/15344843261470602. URL <https://doi.org/10.1177/15344843261470602>.
- Andrew McAfee. The Economic Impact of Generative AI.
- Konstantinos Voudouris, Lucy G. Cheke, and Eric Schulz. Bringing Comparative Cognition To Computers, March 2025. URL <http://arxiv.org/abs/2503.02882>. arXiv:2503.02882 [cs].
- Kaitlyn Zhou, Jena D Hwang, Xiang Ren, Nouha Dziri, Dan Jurafsky, and Maarten Sap. REL-A.I.: An Interaction-Centered Approach To Measuring Human-LM Reliance. *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2025.
- Søren Dinesen Østergaard. Will Generative Artificial Intelligence Chatbots Generate Delusions in Individuals Prone to Psychosis? *Schizophrenia Bulletin*, 49(6):1418–1419, August 2023. ISSN

0586-7614. doi: 10.1093/schbul/sbad128. URL <https://pmc.ncbi.nlm.nih.gov/articles/PMC10686326/>.